

Scalable watermarking for identifying large language model outputs.

Dathathri et al., *Nature* (2024)

Xing Liu

QuantCo reading group

7 Oct 2026

Contents

1. Motivation
2. Background
3. SynthID-Text
4. Guarantees and detection
5. Computational scalability
6. Empirical evidence
7. Critical assessment

Motivation

The Guardian UK

News Opinion Sport Culture Lifestyle

World Europe US news Americas Asia Australia Middle East Africa Inequality Global development

AI (artificial intelligence)

This article is more than 2 months old

AI labels to be compulsory on authentic-looking content under EU rules

Companies must ensure people know when they are interacting with artificially generated images, audio and text designed to look real

Jennifer Rankin in Brussels

Fri 31 Jul 2026 12:31 BST

Share

Prefer the Guardian on Google



Screenshot of an AI-generated image of the French president, Emmanuel Macron, dancing in a 1980s nightclub.

Advertisement

Could you pay less for your pension?

Capital at risk.


[Link.](#) EU AI Act enforcement, starting from 2nd August 2026.

 [Research](#) [Policy](#) [Commitments](#) [Learn](#) [News](#) [Try Claude](#)

Announcements

How Claude's text watermark works

Aug 14, 2026



What is watermarking?

How does watermarking affect Claude's outputs?

Which specific method of watermarking do you use?

What about cases where Claude has profited or misled human users?

Future Claude models will generate text that contains a watermark. This is a way of determining the likelihood that Claude was involved in writing the text, and we, along with several other major AI providers, are implementing this change to comply with the EU AI Act.

In this article, we share answers to some of the questions we've received about how our chosen watermarking method works, whether it affects Claude's outputs, and why we're making this change. To summarize:

[Link.](#) Anthropic adopts text watermarking.

Writing in Acceptance Criteria



Transactions on Machine Learning Research

Follow

2 min read · Aug 28, 2026



9



Three things are true today:

1. Peer reviews generally put least emphasis on the criterion of clear writing. From our experience, this is true for TMLR as well.
2. Journals and conferences including TMLR are seeing a large number of submissions, including potentially AI generated or heavily AI written ones.
3. AI writing is currently often convoluted, has unnecessary jargon, and not easy to understand.

[Link](#). TMLR's call for reviewers to emphasise more on clarity of writing.

Background

How LLMs Generate Texts

- Most LLMs generate texts using autoregressive next-token sampling.



- Example:

	Vocabulary	p_{LM}	Selected token
... my favourite tropical fruit is	mango	0.50	→ lychee
	lychee	0.30	
	papaya	0.15	
	durian	0.05	

Existing Approaches for LLM Text Detection

- Retrieval-based detection
 - ▶ Keeping a growing record of all generated texts and checking against it for matches.
 - ▶ Requires scale and coordination, and raises privacy concerns.
- Post hoc detection
 - ▶ Using the statistical features of text or training a ML-based classifier.
 - ▶ Computationally expensive; inconsistent performance (especially out-of-domain).
- Watermarking
 - ▶ Marking the generated text in some way so that it can subsequently be identified.
 - ▶ Three approaches:
 - ★ During generation. → This paper.
 - ★ Editing already generated text.
 - ★ Altering LLM training data.
 - ▶ Often involves a quality (of generated texts) and effect (of detection) trade-off.

[XL: use plots?]

SynthID-Text

SynthID-Text mechanism




Fig. 1, Dathathri et al.

Guarantees and detection

Computational scalability

Empirical evidence

Critical assessment

